



CONVOLUTIONAL NEURAL NETWORKS FOR IMAGE PROCESSING


Going deeper

Ariel Rokem

Senior Data Scientist, University of Washington



Network with one convolutional layer




Network with one convolutional layer: implementation

```
model = Sequential()  
model.add(Conv2D(10, kernel_size=2, activation='relu',  
                input_shape=(img_rows, img_cols, 1)))  
model.add(Flatten())  
model.add(Dense(3, activation='softmax'))
```



Building a deeper network



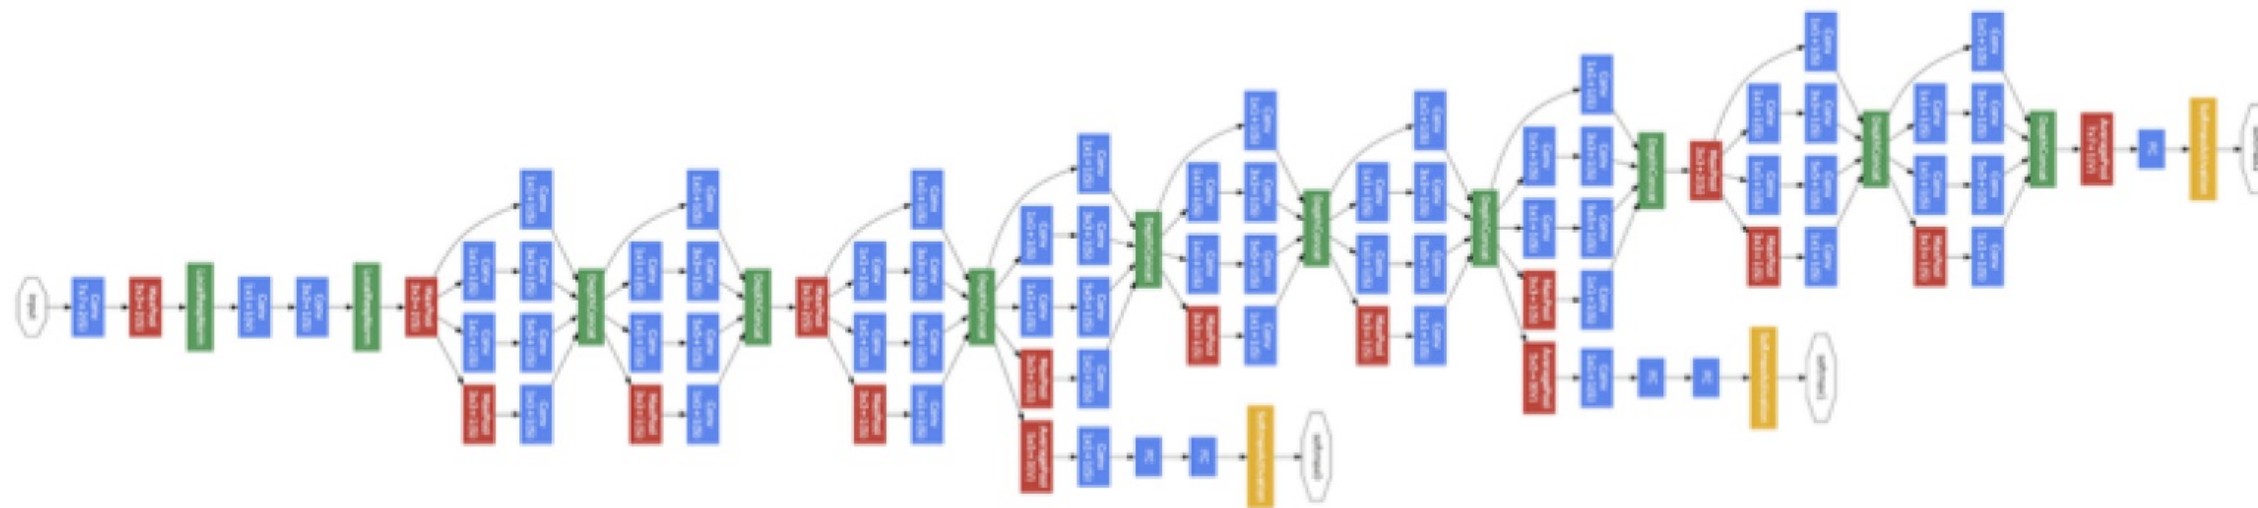
Building a deep network

```
model = Sequential()
model.add(Conv2D(10, kernel_size=2, activation='relu',
                input_shape=(img_rows, img_cols, 1),
                padding='equal'))

# Second convolutional layer
model.add(Conv2D(10, kernel_size=2, activation='relu'))

model.add(Flatten())
model.add(Dense(3, activation='softmax'))
```


Why do we want deep networks?



Convolution
Pooling
Softmax
Other




Features in early layers



Features in intermediate layers



Features in late layers





How deep?

- Depth comes at a computational cost
- May require more data



CONVOLUTIONAL NEURAL NETWORKS FOR IMAGE PROCESSING

Let's practice!



CONVOLUTIONAL NEURAL NETWORKS FOR IMAGE PROCESSING

How many parameters?

Ariel Rokem

Senior Data Scientist, University of Washington



Counting parameters

```
model = Sequential()

model.add(Dense(10, activation='relu',
               input_shape=(784,)))
model.add(Dense(10, activation='relu'))

model.add(Dense(3, activation='softmax'))
```

Model summary

```
# Call the summary method  
model.summary()
```

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 10)	7850
dense_2 (Dense)	(None, 10)	110
dense_3 (Dense)	(None, 3)	33
Total params: 7,993		
Trainable params: 7,993		
Non-trainable params: 0		

Counting parameters

```
model.add(Dense(10,  
                activation='relu',  
                input_shape=(784,)))
```

```
model.add(Dense(10,  
                activation='relu'))
```

```
model.add(Dense(3,  
                activation='softmax'))
```

$$parameters = 784 * 10 + 10$$

$$= 7850$$

$$parameters = 10 * 10 + 10$$

$$= 110$$

$$parameters = 10 * 3 + 3$$

$$= 33$$

$$7850 + 110 + 33 = 7993$$

Model summary

```
model.summary()
```

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 10)	7850
dense_2 (Dense)	(None, 10)	110
dense_3 (Dense)	(None, 3)	33
Total params: 7,993		
Trainable params: 7,993		
Non-trainable params: 0		



The number of parameters in a CNN

```
model = Sequential()
model.add(Conv2D(10, kernel_size=3, activation='relu',
                 input_shape=(28, 28, 1), padding='same'))
model.add(Conv2D(10, kernel_size=3, activation='relu',
                 padding='same'))
model.add(Flatten())
model.add(Dense(3, activation='softmax'))
```

The number of parameters in a CNN

```
model.summary()
```

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 28, 28, 10)	100
conv2d_2 (Conv2D)	(None, 28, 28, 10)	910
flatten_3 (Flatten)	(None, 7840)	0
dense_4 (Dense)	(None, 3)	23523
Total params: 24,533		
Trainable params: 24,533		
Non-trainable params: 0		



The number of parameters in a CNN

```
model.add(Conv2D(10, kernel_size=3,  
                 activation='relu',  
                 input_shape=(28, 28, 1),  
                 padding='same'))
```

```
model.add(Conv2D(10, kernel_size=3,  
                 activation='relu',  
                 padding='same'))
```

```
model.add(Flatten())
```

```
model.add(Dense(3,  
                activation='softmax'))
```

$$parameters = 9 * 10 + 10$$

$$= 100$$

$$parameters = 10 * 9 * 10 + 10$$

$$= 910$$

$$parameters = 0$$

$$parameters = 7840 * 3 + 3$$

$$= 23523$$

$$100 + 910 + 0 + 23523 = 24533$$



Increasing the number of units in each layer

```
model = Sequential()  
model.add(Dense(5, activation='relu',  
                input_shape=(784,), padding='same'))  
model.add(Dense(15, activation='relu', padding='same'))  
model.add(Dense(3, activation='softmax'))
```

Increasing the number of units in each layer

```
model.summary()
```

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 5)	3925
dense_2 (Dense)	(None, 15)	90
dense_3 (Dense)	(None, 3)	48
Total params: 4,063		
Trainable params: 4,063		
Non-trainable params: 0		

Increasing the number of units in each layer

```
model = Sequential()
model.add(Conv2D(5, kernel_size=3, activation='relu',
                 input_shape=(28, 28, 1), padding="same"))
model.add(Conv2D(15, kernel_size=3, activation='relu',
                 padding="same"))
model.add(Flatten())
model.add(Dense(3, activation='softmax'))
```

Increasing the number of units in each layer

```
model.summary()
```

Layer (type)	Output Shape	Param #
conv2d_12 (Conv2D)	(None, 28, 28, 5)	50
conv2d_13 (Conv2D)	(None, 28, 28, 15)	690
flatten_6 (Flatten)	(None, 11760)	0
dense_9 (Dense)	(None, 3)	35283
Total params: 36,023		
Trainable params: 36,023		
Non-trainable params: 0		



CONVOLUTIONAL NEURAL NETWORKS FOR IMAGE PROCESSING

Let's practice!



CONVOLUTIONAL NEURAL NETWORKS FOR IMAGE PROCESSING

Reducing parameters with pooling

Ariel Rokem

Senior Data Scientist, University of Washington


Pooling reduces the size of the output


```
model.summary()
```


Layer (type)	Output Shape	Param #
conv2d_12 (Conv2D)	(None, 28, 28, 5)	50
conv2d_13 (Conv2D)	(None, 28, 28, 15)	690
flatten_6 (Flatten)	(None, 11760)	0
dense_9 (Dense)	(None, 3)	35283
Total params: 36,023		
Trainable params: 36,023		
Non-trainable params: 0		



















Implementing max pooling

```
result = np.zeros((im.shape[0]//2, im.shape[1]//2))  
result[0, 0] = np.max(im[0:2, 0:2])  
result[0, 1] = np.max(im[0:2, 2:4])  
result[0, 2] = np.max(im[0:2, 4:6])
```

...

```
result[1, 0] = np.max(im[2:4, 0:2])  
result[1, 1] = np.max(im[2:4, 2:4])
```

...



Implementing max pooling

```
for ii in range(result.shape[0]):  
    for jj in range(result.shape[1]):  
        result[ii, jj] = np.max(im[ii*2:ii*2+2, jj*2:jj*2+2])
```

Max pooling in Keras

```
from keras.models import Sequential
from keras.layers import Dense, Conv2D, Flatten, MaxPool2D

model = Sequential()
model.add(Conv2D(5, kernel_size=3, activation='relu',
                 input_shape=(img_rows, img_cols, 1)))

model.add(MaxPool2D(2))

model.add(Conv2D(15, kernel_size=3, activation='relu',
                 input_shape=(img_rows, img_cols, 1)))

model.add(MaxPool2D(2))

model.add(Flatten())
model.add(Dense(3, activation='softmax'))
```

Max pooling reduces the number of parameters

```
model.summary()
```

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 26, 26, 5)	50
max_pooling2d_1 (MaxPooling2D)	(None, 13, 13, 5)	0
conv2d_2 (Conv2D)	(None, 11, 11, 15)	690
max_pooling2d_2 (MaxPooling2D)	(None, 5, 5, 15)	0
flatten_1 (Flatten)	(None, 375)	0
dense_1 (Dense)	(None, 3)	1128
Total params: 1,868		
Trainable params: 1,868		
Non-trainable params: 0		



CONVOLUTIONAL NEURAL NETWORKS FOR IMAGE PROCESSING

Let's practice!